



MONASH
University

MONASH
BUSINESS
SCHOOL

Tidy forecasting in R



Rob J Hyndman

26 September 2019

Outline

- 1 Tidy time series data
- 2 Benchmark forecasting methods
- 3 Exponential smoothing
- 4 ARIMA models
- 5 Forecast accuracy measures

Outline

- 1 Tidy time series data
- 2 Benchmark forecasting methods
- 3 Exponential smoothing
- 4 ARIMA models
- 5 Forecast accuracy measures



tsibble



tsibbledata



feasts



Sable

Time series data

- Four-yearly Olympic winning times
- Annual Google profits
- Quarterly Australian beer production
- Monthly rainfall
- Weekly retail sales
- Daily IBM stock prices
- Hourly electricity demand
- 5-minute freeway traffic counts
- Time-stamped stock transaction data

tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
## # Key:      Country [263]
##   Year Country      GDP Imports Exports Population
##   <dbl> <fct>      <dbl>   <dbl>   <dbl>      <dbl>
## 1  1960 Afghanistan 5377777811.    7.02    4.13    8996351
## 2  1961 Afghanistan 5488888896.    8.10    4.45    9166764
## 3  1962 Afghanistan 5466666678.    9.35    4.88    9345868
## 4  1963 Afghanistan 7511111191.   16.9     9.17    9533954
## 5  1964 Afghanistan 8000000044.   18.1     8.89    9731361
## 6  1965 Afghanistan 10066666638.  21.4    11.3    9938414
## 7  1966 Afghanistan 13999999967.  18.6     8.57   10152331
## 8  1967 Afghanistan 16733333418.  14.2     6.77   10372630
## 9  1968 Afghanistan 13733333367.  15.2     8.90   10604346
## 10 1969 Afghanistan 14088888922.  15.0    10.1   10854428
## # ... with 15,140 more rows
```

tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
## # Key:      Country [263]
##   Year Country      GDP Imports Exports Population
##   Index <fct>      <dbl>   <dbl>   <dbl>         <dbl>
## 1 1960 Afghanistan 5377777811.    7.02    4.13    8996351
## 2 1961 Afghanistan 5488888896.    8.10    4.45    9166764
## 3 1962 Afghanistan 5466666678.    9.35    4.88    9345868
## 4 1963 Afghanistan 7511111191.   16.9     9.17    9533954
## 5 1964 Afghanistan 8000000044.   18.1     8.89    9731361
## 6 1965 Afghanistan 10066666638.  21.4    11.3    9938414
## 7 1966 Afghanistan 13999999967.  18.6     8.57   10152331
## 8 1967 Afghanistan 16733333418.  14.2     6.77   10372630
## 9 1968 Afghanistan 13733333367.  15.2     8.90   10604346
## 10 1969 Afghanistan 14088888922.  15.0    10.1   10854428
## # ... with 15,140 more rows
```

tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
```

```
## # Key:          Country [263]
```

```
##      Year Country      GDP Imports Exports Population
##      Index  Key      <dbl>   <dbl>   <dbl>         <dbl>
##  1  1960 Afghanistan 5377777811.    7.02    4.13    8996351
##  2  1961 Afghanistan 5488888896.    8.10    4.45    9166764
##  3  1962 Afghanistan 5466666678.    9.35    4.88    9345868
##  4  1963 Afghanistan 7511111191.   16.9    9.17    9533954
##  5  1964 Afghanistan 8000000044.   18.1    8.89    9731361
##  6  1965 Afghanistan 10066666638.   21.4   11.3    9938414
##  7  1966 Afghanistan 13999999967.   18.6    8.57   10152331
##  8  1967 Afghanistan 16733333418.   14.2    6.77   10372630
##  9  1968 Afghanistan 13733333367.   15.2    8.90   10604346
## 10  1969 Afghanistan 14088888922.   15.0   10.1   10854428
## # ... with 15,140 more rows
```


tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
```

```
## # Key:      Country [263]
```

```
##      Year Country      GDP Imports Exports Population
```

```
##      Index  Key      Measured variables
```

```
## 1  1960 Afghanistan 5377777811.    7.02    4.13    8996351
```

```
## 2  1961 Afghanistan 5488888896.    8.10    4.45    9166764
```

```
## 3  1962 Afghanistan 5466666678.    9.35    4.88    9345868
```

```
## 4  1963 Afghanistan 7511111191.   16.9    9.17    9533954
```

```
## 5  1964 Afghanistan 8000000044.   18.1    8.89    9731361
```

```
## 6  1965 Afghanistan 10066666638.  21.4   11.3    9938414
```

```
## 7  1966 Afghanistan 13999999967.  18.6    8.57   10152331
```

```
## 8  1967 Afghanistan 16733333418.  14.2    6.77   10372630
```

```
## 9  1968 Afghanistan 13733333367.  15.2    8.90   10604346
```

```
## 10 1969 Afghanistan 14088888922.  15.0   10.1   10854428
```

```
## # ... with 15,140 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:           Region, State, Purpose [304]
##   Quarter Region   State Purpose   Trips
##   <qtr> <chr>      <chr> <chr>    <dbl>
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:           Region, State, Purpose [304]
##   Quarter Region  State Purpose  Trips
##   Index  <chr>    <chr> <chr>    <dbl>
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:      Region, State, Purpose [304]
##   Quarter Region  State Purpose  Trips
##   Index      Keys      <dbl>
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
```

```
## # Key:           Region, State, Purpose [304]
```

```
##   Quarter Region State Purpose Trips
```

```
##   Index      Keys      Measure
```

```
## 1 1998 Q1 Adelaide SA      Business 135.
```

```
## 2 1998 Q2 Adelaide SA      Business 110.
```

```
## 3 1998 Q3 Adelaide SA      Business 166.
```

```
## 4 1998 Q4 Adelaide SA      Business 127.
```

```
## 5 1999 Q1 Adelaide SA      Business 137.
```

```
## 6 1999 Q2 Adelaide SA      Business 200.
```

```
## 7 1999 Q3 Adelaide SA      Business 169.
```

```
## 8 1999 Q4 Adelaide SA      Business 134.
```

```
## 9 2000 Q1 Adelaide SA      Business 154.
```

```
## 10 2000 Q2 Adelaide SA      Business 169.
```

```
## # ... with 24,310 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:      Region, State, Purpose [304]
##   Quarter Region  State Purpose  Trips
##   Index      Keys      Measure
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```

Domestic visitor
nights in thousands
by state/region and
purpose.


tsibble objects

- A `tsibble` allows storage and manipulation of multiple time series in R.
- It contains:
 - ▶ An index: time information about the observation
 - ▶ Measured variable(s): numbers of interest
 - ▶ Key variable(s): optional unique identifiers for each series
- It works with tidyverse functions.

Outline

- 1 Tidy time series data
- 2 Benchmark forecasting methods
- 3 Exponential smoothing
- 4 ARIMA models
- 5 Forecast accuracy measures

Benchmark forecasting methods




How would you forecast these series?

Benchmark forecasting methods

Mean method


- Forecast of all future values is equal to mean of historical data $\{y_1, \dots, y_T\}$.
- Forecasts: $\hat{y}_{T+h|T} = \bar{y} = (y_1 + \dots + y_T)/T$



Benchmark forecasting methods

Naïve method


- Forecasts equal to last observed value.
- Forecasts: $\hat{y}_{T+h|T} = y_T$.
- Consequence of efficient market hypothesis.



Benchmark forecasting methods

Seasonal naïve method


- Forecasts equal to last value from same season.
- Forecasts: $\hat{y}_{T+h|T} = y_{T+h-m(k+1)}$, where m = seasonal period and k is the integer part of $(h - 1)/m$.



Benchmark forecasting methods


Drift method

- Forecasts equal to last value plus average change.
- Forecasts: $\hat{y}_{T+h|T} = y_T + \frac{h}{T-1}(y_T - y_1)$.
- Equivalent to line between first and last observations.



Data preparation and visualisation

```
# Set training data from 1970 to 2006  
train <- aus_production %>%  
  filter(between(year(Quarter), 1970, 2006))  
train %>% autoplot(Bricks)
```



Model estimation

The `model()` function trains models to data.

```
# Fit the models
fit <- train %>%
  model(
    Mean = MEAN(Bricks),
    Naïve = NAIVE(Bricks),
    SeasonalNaïve = SNAIVE(Bricks),
    Drift = RW(Bricks ~ drift())
  )
```

```
fit
```

```
## # A mable: 1 x 4
##   Mean      Naïve    SeasonalNaïve  Drift
##   <model> <model> <model>          <model>
## 1 <MEAN>  <NAIVE> <SNAIVE>          <RW w/ drift>
```

Producing forecasts

```
fc <- fit %>%  
  forecast(h = 11)
```


```
## # A fable: 44 x 4 [1Q]  
## # Key:      .model [4]  
##   .model Quarter Bricks .distribution  
##   <chr>      <qtr>    <dbl> <dist>  
## 1 Mean      2007 Q1     450. N(450, 4030)  
## 2 Mean      2007 Q2     450. N(450, 4030)  
## 3 Mean      2007 Q3     450. N(450, 4030)  
## 4 Mean      2007 Q4     450. N(450, 4030)  
## # ... with 40 more rows
```

A fable is a forecast table with point forecasts and distributions.

Visualising forecasts

```
fc %>%
```

```
  autoplot(train, level = NULL) +  
  ggtitle("Forecasts for quarterly brick production") +  
  xlab("Year") +  
  guides(colour=guide_legend(title="Forecast"))
```



Forecasting many series

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:      Region, State, Purpose [304]
##   Quarter Region   State Purpose   Trips
##   <qtr> <chr>      <chr> <chr>      <dbl>
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
```

Forecasting many series

```
tourism %>%  
  model(  
    mean = MEAN(Trips),  
    snaive = SNAIVE(Trips)  
  )
```

```
## # A mable: 304 x 5
```

```
## # Key:      Region, State, Purpose [304]
```

	Region	State	Purpose	mean	snaive
	<chr>	<chr>	<chr>	<model>	<model>
## 1	Adelaide	SA	Business	<MEAN>	<SNAIVE>
## 2	Adelaide	SA	Holiday	<MEAN>	<SNAIVE>
## 3	Adelaide	SA	Other	<MEAN>	<SNAIVE>
## 4	Adelaide	SA	Visiting	<MEAN>	<SNAIVE>
## 5	Adelaide Hills	SA	Business	<MEAN>	<SNAIVE>
## 6	Adelaide Hills	SA	Holiday	<MEAN>	<SNAIVE>

Forecasting many series

```
tourism %>%  
  model(  
    mean = MEAN(Trips),  
    snaive = SNAIVE(Trips)  
  ) %>%  
  forecast(h= "2 years")
```

```
## # A fable: 4,864 x 7 [1Q]  
## # Key:      Region, State, Purpose, .model [608]  
##   Region State Purpose .model   Quarter Trips  
##   <chr>  <chr> <chr>   <chr>      <qtr> <dbl>  
## 1 Adela~ SA     Busine~ mean    2018 Q1  156.  
## 2 Adela~ SA     Busine~ mean    2018 Q2  156.  
## 3 Adela~ SA     Busine~ mean    2018 Q3  156.  
## 4 Adela~ SA     Busine~ mean    2018 Q4  156.  
## 5 Adela~ SA     Busine~ mean    2019 Q1  156.
```

Outline

- 1 Tidy time series data
- 2 Benchmark forecasting methods
- 3 Exponential smoothing
- 4 ARIMA models
- 5 Forecast accuracy measures

Historical perspective

- Developed in the 1950s and 1960s as methods (algorithms) to produce point forecasts.
- Combine a “level”, “trend” (slope) and “seasonal” component to describe a time series.
- The rate of change of the components are controlled by “smoothing parameters”: α , β and γ respectively.
- Need to choose best values for the smoothing parameters (and initial states).
- Equivalent ETS state space models developed in the 1990s and 2000s.

ETS models

General notation

ETS : ExponenTial Smoothing



Error Trend Season

The diagram shows three arrows pointing upwards from the words 'Error', 'Trend', and 'Season' to the letters 'E', 'T', and 'S' respectively in the 'ETS' part of the text above.

Error: Additive ("A") or multiplicative ("M")

ETS models

General notation

ETS : ExponenTial Smoothing

Error Trend Season


Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

ETS models

General notation

ETS : ExponenTial Smoothing
Error Trend Season



Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

Seasonality: None ("N"), additive ("A") or multiplicative ("M")

Exponential smoothing models

Additive Error

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
Trend Component	N (None)	A,N,N	A,N,A	A,N,M
	A (Additive)	A,A,N	A,A,A	A,A,M
	A _d (Additive damped)	A,A _d ,N	A,A _d ,A	A,A_d,M

Multiplicative Error

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
Trend Component	N (None)	M,N,N	M,N,A	M,N,M
	A (Additive)	M,A,N	M,A,A	M,A,M
	A _d (Additive damped)	M,A _d ,N	M,A _d ,A	M,A _d ,M

Automatic forecasting

From Hyndman et al. (IJF, 2002):

- 1 Apply each model that is appropriate to the data. Optimize parameters and initial values using MLE.
 - 2 Select best method using AICc.
 - 3 Produce forecasts using best method.
 - 4 Obtain forecast intervals using underlying state space model.
- Method performed very well in M3 competition.
 - Used as a benchmark in the M4 competition.

Example: Australian population

```
aus_economy <- global_economy %>% filter(Country == "Australia") %>%  
  mutate(Pop = Population/1e6)  
fit <- aus_economy %>% model(best = ETS(Pop))  
report(fit)
```

```
## Series: Pop  
## Model: ETS(A,A,N)  
## Smoothing parameters:  
##   alpha = 1  
##   beta  = 0.327  
##  
## Initial states:  
##   l      b  
## 10.1 0.222  
##  
## sigma^2: 0.0041  
##  
## AIC AICc BIC  
## -77.0 -75.8 -66.7
```


Example: Australian population

```
components(fit)
```

```
## # A dable:                59 x 7 [1Y]
## # Key:                    Country, .model [1]
## # ETS(A,A,N) Decomposition: Pop = lag(level, 1) +
## #   lag(slope, 1) + remainder
##   Country   .model Year   Pop level slope remainder
##   <fct>     <chr> <dbl> <dbl> <dbl> <dbl>      <dbl>
## 1 Australia best   1959  NA    10.1 0.222 NA
## 2 Australia best   1960  10.3  10.3 0.222 -0.000145
## 3 Australia best   1961  10.5  10.5 0.217 -0.0159
## 4 Australia best   1962  10.7  10.7 0.231  0.0418
## 5 Australia best   1963  11.0  11.0 0.223 -0.0229
## 6 Australia best   1964  11.2  11.2 0.221 -0.00641
## 7 Australia best   1965  11.4  11.4 0.221 -0.000314
## 8 Australia best   1966  11.7  11.7 0.235  0.0418
## 9 Australia best   1967  11.8  11.8 0.206 -0.0869
## 10 Australia best  1968  12.0  12.0 0.208  0.00350
```


Example: Australian population

```
components(fit) %>% autoplot()
```



Example: Australian population

```
fit %>%  
  forecast(h = 20) %>%  
  autoplot(aus_economy) +  
  ylab("Population") + xlab("Year")
```



Example: National populations

```
fit <- global_economy %>%  
  mutate(Pop = Population/1e6) %>%  
  model(ets = ETS(Pop))  
fit
```

```
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country      ets  
##   <fct>        <model>  
## 1 Afghanistan <ETS(A,A,N)>  
## 2 Albania      <ETS(M,A,N)>  
## 3 Algeria      <ETS(M,A,N)>  
## 4 American Samoa <ETS(M,A,N)>  
## 5 Andorra      <ETS(M,A,N)>  
## 6 Angola        <ETS(M,A,N)>  
## 7 Antigua and Barbuda <ETS(M,A,N)>  
## 8 Arab World   <ETS(M,A,N)>  
## 9 Argentina    <ETS(A,A,N)>  
## 10 Armenia     <ETS(M,A,N)>  
## # ... with 253 more rows
```


Example: National populations

```
fit %>%  
  forecast(h = 5)
```

```
## # A tibble: 1,315 x 5 [1Y]  
## # Key:      Country, .model [263]  
##   Country      .model Year   Pop .distribution  
##   <fct>        <chr>   <dbl> <dbl> <dist>  
## 1 Afghanistan ets      2018  36.4 N(36, 0.012)  
## 2 Afghanistan ets      2019  37.3 N(37, 0.059)  
## 3 Afghanistan ets      2020  38.2 N(38, 0.164)  
## 4 Afghanistan ets      2021  39.0 N(39, 0.351)  
## 5 Afghanistan ets      2022  39.9 N(40, 0.644)  
## 6 Albania     ets      2018   2.87 N(2.9, 0.00012)  
## 7 Albania     ets      2019   2.87 N(2.9, 0.00060)  
## 8 Albania     ets      2020   2.87 N(2.9, 0.00169)  
## 9 Albania     ets      2021   2.86 N(2.9, 0.00362)
```

Example: Australian holiday tourism

```
holidays <- tourism %>%  
  filter(Purpose == "Holiday")  
fit <- holidays %>% model(ets = ETS(Trips))  
fit
```

```
## # A mable: 76 x 4  
## # Key:      Region, State, Purpose [76]  
##   Region                State Purpose ets  
##   <chr>                 <chr> <chr>  <model>  
## 1 Adelaide             SA     Holiday <ETS(A,N,A)>  
## 2 Adelaide Hills       SA     Holiday <ETS(A,A,N)>  
## 3 Alice Springs        NT     Holiday <ETS(M,N,A)>  
## 4 Australia's Coral Coast WA     Holiday <ETS(M,N,A)>  
## 5 Australia's Golden Outback WA     Holiday <ETS(M,N,M)>  
## 6 Australia's North West WA     Holiday <ETS(A,N,A)>  
## 7 Australia's South West WA     Holiday <ETS(M,N,M)>  
## 8 Ballarat             VIC     Holiday <ETS(M,N,A)>  
## 9 Barkly               NT     Holiday <ETS(A,N,A)>  
## 10 Barossa             SA     Holiday <ETS(A,N,N)>
```

Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% report()
```

```
## Series: Trips
## Model: ETS(M,N,A)
## Smoothing parameters:
##   alpha = 0.157
##   gamma = 1e-04
##
## Initial states:
##   l   s1  s2    s3    s4
## 142 -61 131 -42.2 -27.7
##
##   sigma^2: 0.0388
##
## AIC AICc BIC
## 852 854 869
```

Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% components(fit)
```


```
## # A dable:                84 x 9 [1Q]
## # Key:                    Region, State, Purpose, .model
## #   [1]
## # ETS(M,N,A) Decomposition: Trips = (lag(level, 1) +
## #   lag(season, 4)) * (1 + remainder)
##   Region State Purpose .model   Quarter Trips level season
##   <chr>   <chr> <chr>   <chr>       <qtr> <dbl>  <dbl>  <dbl>
## 1 Snowy~ NSW   Holiday ets       1997 Q1   NA      NA    -27.7
## 2 Snowy~ NSW   Holiday ets       1997 Q2   NA      NA    -42.2
## 3 Snowy~ NSW   Holiday ets       1997 Q3   NA      NA    131.
## 4 Snowy~ NSW   Holiday ets       1997 Q4   NA     142.   -61.0
## 5 Snowy~ NSW   Holiday ets       1998 Q1  101.    140.   -27.7
## 6 Snowy~ NSW   Holiday ets       1998 Q2  112.    142.   -42.2
## 7 Snowy~ NSW   Holiday ets       1998 Q3  310.    148.    131.
## 8 Snowy~ NSW   Holiday ets       1998 Q4   89.8   148.   -61.0
## 9 Snowy~ NSW   Holiday ets       1999 Q1  112.    147.   -27.7
## 10 Snowy~ NSW   Holiday ets       1999 Q2  103.    147.   -42.2
```

Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>%  
  components(fit) %>% autoplot()
```

ETS(M,N,A) decomposition

Trips = (lag(level, 1) + lag(season, 4)) * (1 + remainder)




Example: Australian holiday tourism

```
fit %>% forecast()
```

```
## # A tibble: 608 x 7 [1Q]
## # Key:   Region, State, Purpose, .model [76]
##   Region State Purpose .model   Quarter Trips
##   <chr>  <chr> <chr>   <chr>         <qtr> <dbl>
## 1 Adela~ SA    Holiday ets        2018 Q1  210.
## 2 Adela~ SA    Holiday ets        2018 Q2  173.
## 3 Adela~ SA    Holiday ets        2018 Q3  169.
## 4 Adela~ SA    Holiday ets        2018 Q4  186.
## 5 Adela~ SA    Holiday ets        2019 Q1  210.
## 6 Adela~ SA    Holiday ets        2019 Q2  173.
## 7 Adela~ SA    Holiday ets        2019 Q3  169.
## 8 Adela~ SA    Holiday ets        2019 Q4  186.
## 9 Adela~ SA    Holiday ets        2018 Q1   19.4
## 10 Adela~ SA    Holiday ets        2018 Q2   19.6
## # ... with 598 more rows, and 1 more variable:
## #   .distribution <dist>
```

Example: Australian holiday tourism

```
fit %>% forecast() %>%  
  filter(Region=="Snowy Mountains") %>%  
  autoplot(holidays) +  
    xlab("Year") + ylab("Overnight trips (thousands)")
```



Outline

- 1 Tidy time series data
- 2 Benchmark forecasting methods
- 3 Exponential smoothing
- 4 ARIMA models
- 5 Forecast accuracy measures

ARIMA models

- AR:** autoregressive (lagged observations as inputs)
- I:** integrated (differencing to make series stationary)
- MA:** moving average (lagged errors as inputs)

ARIMA models

AR: autoregressive (lagged observations as inputs)

I: integrated (differencing to make series stationary)

MA: moving average (lagged errors as inputs)

An ARIMA model is rarely interpretable in terms of visible data structures like trend and seasonality. But it can capture a huge range of time series patterns.

ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of y_t** and **lagged errors**.

ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of y_t** and **lagged errors**.

Autoregressive Integrated Moving Average models

- Combine ARMA model with **differencing**.
- d -differenced series follows an ARMA model.
- Need to choose p, d, q and whether or not to include c .

Seasonal ARIMA models

ARIMA	$\underbrace{(p, d, q)}$	$\underbrace{(P, D, Q)_m}$
	↑	↑
	Non-seasonal part of the model	Seasonal part of of the model

- m = number of observations per year.
- d first differences, D seasonal differences
- p AR lags, q MA lags
- P seasonal AR lags, Q seasonal MA lags

Seasonal and non-seasonal terms combine multiplicatively

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p , q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

Note: Can't compare AICc for different values of d .

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

How does ARIMA() work?

Step 1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

Step 2: Consider variations of current model:

- vary one of p , q , from current model by ± 1 ;
- p , q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

Step 2: Consider variations of current model:

- vary one of p , q , from current model by ± 1 ;
- p , q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

Repeat Step 2 until no lower AICc can be found.


Example: National populations

```
fit <- global_economy %>%  
  model(arima = ARIMA(Population))  
fit
```

```
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country      arima  
##   <fct>        <model>  
## 1 Afghanistan <ARIMA(4,2,1)>  
## 2 Albania     <ARIMA(0,2,2)>  
## 3 Algeria     <ARIMA(2,2,2)>  
## 4 American Samoa <ARIMA(2,2,2)>  
## 5 Andorra     <ARIMA(2,1,2) w/ drift>  
## 6 Angola      <ARIMA(4,2,1)>  
## 7 Antigua and Barbuda <ARIMA(2,1,2) w/ drift>  
## 8 Arab World  <ARIMA(0,2,1)>
```

Example: National populations

```
fit %>% forecast(h=10) %>%  
  filter(Country=="Australia") %>%  
  autoplot(global_economy)
```



Outline

- 1 Tidy time series data
- 2 Benchmark forecasting methods
- 3 Exponential smoothing
- 4 ARIMA models
- 5 Forecast accuracy measures

Training and test sets



- A model which fits the training data well will not necessarily forecast well.
- Forecast accuracy is based only on the test set.


Forecast errors

Forecast “error”: the difference between an observed value and its forecast.

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T},$$

where the training data is given by $\{y_1, \dots, y_T\}$

Measures of forecast accuracy



Measures of forecast accuracy

y_{T+h} = $(T + h)$ th observation, $h = 1, \dots, H$

$\hat{y}_{T+h|T}$ = its forecast based on data up to time T .

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$$

$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

Measures of forecast accuracy

y_{T+h} = $(T + h)$ th observation, $h = 1, \dots, H$

$\hat{y}_{T+h|T}$ = its forecast based on data up to time T .

$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$

MAE = $\text{mean}(|e_{T+h}|)$

MSE = $\text{mean}(e_{T+h}^2)$

RMSE = $\sqrt{\text{mean}(e_{T+h}^2)}$

MAPE = $100\text{mean}(|e_{T+h}|/|y_{T+h}|)$

- MAE, MSE, RMSE are all scale dependent.
- MAPE is scale independent but is only sensible if $y_t \gg 0$ for all t , and y has a natural zero.

Measures of forecast accuracy

Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|/Q)$$

where Q is a stable measure of the scale of the time series $\{y_t\}$.

Proposed by Hyndman and Koehler (IJF, 2006).

For non-seasonal time series,

$$Q = (T - 1)^{-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

works well. Then MASE is equivalent to MAE relative to a naïve method.

Measures of forecast accuracy

Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|/Q)$$

where Q is a stable measure of the scale of the time series $\{y_t\}$.

Proposed by Hyndman and Koehler (IJF, 2006).

For seasonal time series,

$$Q = (T - m)^{-1} \sum_{t=m+1}^T |y_t - y_{t-m}|$$

works well. Then MASE is equivalent to MAE relative to a seasonal naïve method.

Measures of forecast accuracy

```
recent_production <- aus_production %>%  
  filter(year(Quarter) >= 1992)  
train <- recent_production %>% filter(year(Quarter) <= 2007)  
beer_fit <- train %>%  
  model(  
    ets = ETS(Beer),  
    arima = ARIMA(Beer)  
  )  
beer_fc <- forecast(beer_fit, h="4 years")  
accuracy(beer_fc, aus_production)
```

```
## # A tibble: 2 x 9  
##   .model .type    ME  RMSE   MAE   MPE  MAPE  MASE  ACF1  
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 arima  Test   4.18  11.2  10.4  0.940  2.47  0.657  0.145  
## 2 ets    Test   0.854  9.80  8.99  0.151  2.18  0.568  0.207
```